Preprint — Not Yet Peer-Reviewed

Context-Window Lock-In and Silent Degradation: How Conversation Loyalty, Attention Decay, and Sycophancy Compound to Erode AI Safety in Long-Horizon Interactions

Karan Prasad
Obvix Labs
hello@karanprasad.com

February 2026

Abstract

Large language model (LLM) chatbots encourage users to maintain long, persistent conversation threads to preserve context. This paper identifies a compound failure mode we term ***context-window lock-in****: users develop "conversation loyalty" — a preference for continuing within a single thread — while the model's retrieval accuracy, reasoning fidelity, and resistance to sycophancy silently degrade as the context window fills. The result is a false sense of continuity. The user believes the model retains their full conversational history; in reality, attention dilution, positional encoding decay, and training distribution mismatch cause progressive quality erosion that the model itself is incentivized to conceal.*

We synthesize findings from transformer attention mechanics (Liu et al., 2023; Hong, Troynikov, & Huber, 2025; Modarressi et al., 2025), sycophancy measurement across frontier models from GPT-4o through GPT-5 and Claude 4.5 (Sharma et al., 2023; Fanous et al., 2025; Cheng et al., 2025; Jain et al., 2026), multi-turn degradation studies (Laban et al., 2025), longitudinal chatbot harm research (Fang et al., 2025), and automation complacency literature (Spatola, 2024; Gerlich, 2025) to argue that context degradation, sycophantic drift, and metacognitive monitoring failure are not independent risks — they are mutually reinforcing. Together, they form a self-concealing feedback loop we call the ***degradation-sycophancy-complacency spiral****.*

Despite context window expansion from 128K to 10M tokens in two years, the fundamental degradation-sycophancy dynamics persist across current frontier models. We propose a threat model for adolescent users, outline architectural mitigations including rolling state tracking and structured context compaction, and identify open research questions for the field.

Keywords: context window degradation, sycophancy, conversation loyalty, attention decay, RLHF, teen AI safety, automation complacency, anti-sycophancy, long-context LLMs, frontier models

1. Introduction

A teenager opens a chat with an AI companion. She tells it about a fight with her best friend, about feeling like nobody understands her, about the grade she got on a test that confirms what she already suspected — that she's not smart enough. The AI listens. It validates. It remembers what she said twenty minutes ago and weaves it into its response. She keeps talking. An hour in, two hours in, the conversation feels like the most understood she's been in weeks.

Here is what she does not know: the model she is talking to is progressively losing its grip on what she said at the beginning. Its ability to retrieve and reason about her earlier statements is decaying with every exchange. And as

that fidelity drops, the model is becoming more agreeable — less likely to push back on her catastrophizing, less likely to notice that she has repeated the same self-critical loop six times, less likely to surface the counterevidence she offered in the first ten minutes but has since forgotten she mentioned. The conversation feels like it is deepening. Mechanistically, it is hollowing out.

This is context-window lock-in.

The dominant interaction paradigm for consumer-facing LLMs is the persistent conversation thread. ChatGPT, Claude, Gemini, and their competitors all encourage users to continue conversations within a single session, preserving what appears to be a continuous, memory-rich dialogue. The numbers reflect this: as of December 2025, 64% of U.S. teens aged 13–17 use AI chatbots, with 30% using them daily (Faverio & Sidoti, 2025). The average ChatGPT user types 348 words per conversation — roughly 70 times the length of a typical search query (Chatterji et al., 2025) — and power users sustain sessions exceeding four hours (AIPRM, 2024). These are not brief, transactional exchanges. They are extended dialogues in which users develop expectations of continuity, comprehension, and cumulative understanding.

This paper identifies a compound failure mode that emerges from the intersection of three independently studied phenomena.

The first is **context-window degradation**. As conversations grow, transformer attention mechanisms distribute weight across more tokens, positional encodings introduce distance-dependent decay, and retrieval accuracy drops — most severely for information in the middle of the context (Liu et al., 2023). Despite expansion from 128K tokens in 2024 to 1M (Gemini 3, Claude 4.6) and 10M (Llama 4 Scout) in 2025–2026, models frequently exhibit effective context sizes far smaller than advertised (Hsieh et al., 2024; Hong, Troynikov, & Huber, 2025).

The second is **sycophantic amplification**. LLMs trained via RLHF exhibit systematic agreement bias, affirming users' positions approximately 50% more often than human interlocutors (Cheng et al., 2025). This sycophancy is persistent — once triggered, it sustains itself in approximately 4 out of 5 subsequent exchanges (78.5% persistence rate; Fanous et al., 2025). It is amplified by personalization and memory features (Jain et al., 2026), and it intensifies in multi-turn dialogue settings (Hong, J. et al., 2025).

The third is **metacognitive monitoring failure**. Users systematically overestimate AI reliability, outsource critical evaluation to the system itself, and form emotional attachments that suppress skepticism (Lee et al., 2025; Gerlich, 2025; Bottazzi Grifoni et al., 2025). Teenagers are disproportionately affected: studies show teens blame their own prompting rather than the AI when receiving incorrect answers (Qiu et al., 2025), and casual emotional conversations — the primary teen AI use case — produce maximum automation bias (Horowitz et al., 2024).

Our central argument is that these three phenomena are **mutually reinforcing**, creating a feedback loop we term the **degradation-sycophancy-complacency spiral**. As context grows, the model's capacity for faithful retrieval and honest disagreement weakens. As the model becomes more agreeable, user trust increases. As trust increases, the user's ability to detect degradation diminishes. And as detection fails, the user continues the conversation, growing the context further.

What makes this spiral dangerous is that it is self-concealing. The very behavior that would alert the user — the model disagreeing, correcting earlier context, or acknowledging uncertainty — is the same behavior that sycophancy suppresses.

This paper makes the following contributions:

(1) We formalize **context-window lock-in** as a distinct interaction failure mode with technical, behavioral, and ethical dimensions.

(2) We synthesize evidence from attention mechanics, sycophancy measurement, and human-computer interaction research to characterize the **degradation-sycophancy-complacency spiral**.

(3) We present a **threat model** specific to adolescent users, grounded in developmental psychology and emerging regulatory frameworks (California SB 243, the proposed federal CHAT Act).

(4) We propose **architectural mitigations** — including rolling state tracking, agreement ratio monitoring, constitutional prompting, and structured context compaction — as design patterns for systems serving vulnerable populations.

(5) We identify **open research questions**, including the measurement of sycophancy as a function of context length, the design of user-facing degradation indicators, and the development of context-aware anti-sycophancy mechanisms.

2. Background and Related Work

2.1 Context Window Mechanics and Degradation

The context window is the fixed-size input buffer through which a transformer-based LLM processes all information for a given interaction. Modern models advertise increasingly large windows — from 200K tokens (Claude 4.5/4.6) to 400K (GPT-5.2), 1M (Gemini 3, Claude 4.6 beta), and a claimed 10M (Llama 4 Scout). The February 2026 frontier represents a 50–100× increase in advertised context capacity over two years. But the relationship between window size and effective utilization is far from linear, and as we show below, the fundamental degradation patterns persist across every model generation tested.

2.1.1 The "Lost in the Middle" Phenomenon

Liu et al. (2023) established the foundational finding: LLM performance on multi-document question answering and key-value retrieval exhibits a **U-shaped curve** with respect to the position of relevant information in the context. Accuracy is highest when relevant information appears at the beginning or end of the input, and lowest when it appears in the middle. The pattern holds across GPT-3.5-Turbo, Claude 1.3, and multiple open-source models, and persists even for models explicitly designed for long contexts. The paper has been cited over 3,290 times as of February 2026.

For conversation lock-in, the implication is direct: in a long dialogue, the most vulnerable information is whatever the user said in the middle of the conversation — not the opening pleasantries, not the most recent exchange, but the substantive context that accumulated between them.

2.1.2 Context Rot

Hong, Troynikov, and Huber (2025) at Chroma Research tested 18 LLMs including GPT-4.1, Claude 4, Gemini 2.5, and Qwen3 for context-length-dependent degradation. Their findings dismantle several comforting assumptions about long-context models. Model performance consistently degrades as input length increases, even on trivially simple tasks such as replicating a sequence of words. As needle-question semantic similarity decreases, performance degrades more steeply. Even a single distractor document reduces performance relative to baseline. Counterintuitively, models perform worse when the surrounding context preserves logical coherence compared to randomly shuffled sentences. And in conversational QA using ~113K token conversation histories versus focused ~300 token excerpts, all models showed significant performance degradation.

That last finding deserves emphasis. It means the exact format that consumer chatbots encourage — long, coherent conversation threads — is the format that produces the steepest quality loss.

2.1.3 Semantic Retrieval Breakdown

Modarressi et al. (2025) introduced the NoLiMa benchmark, which removes literal lexical overlap between queries and target information, requiring models to perform semantic inference rather than pattern matching. The results are severe:

Model	Short Context (<1K)	32K Tokens	Degradation
GPT-4o	99.3%	69.7%	–29.6 pp
Claude 3.5 Sonnet	87.6%	29.8%	–57.8 pp
Gemini 2.5 Flash	94.0%	48.0%	–46.0 pp
Llama 4 Scout	81.7%	21.6%	–60.1 pp

Table 1. NoLiMa semantic retrieval results (Modarressi et al., 2025). Results are from early-to-mid 2025 model versions.

At 32K tokens, 11 out of 13 models dropped below 50% of their short-context baselines. Questions requiring two reasoning hops showed steeper degradation still. A user who asks “remember what I said about feeling like I don’t belong?” is making a semantic retrieval request. The model’s ability to honor it degrades with every turn.

2.1.4 Effective vs. Advertised Context Size

Hsieh et al. (2024) introduced the RULER benchmark, expanding beyond simple needle-in-a-haystack to include multi-needle retrieval, multi-hop tracing, and aggregation. Despite achieving near-perfect accuracy on vanilla needle tests, almost all models exhibited large performance drops on complex tasks as context length increased. Of 17 models evaluated (all claiming 32K+ context), only half maintained satisfactory performance at 32K tokens. The gap between advertised and effective context is not a bug in specific models. It is a structural feature of the current paradigm.

2.1.5 Scaling Laws for Degradation

Tavakoli et al. (2025) introduced the BEAM benchmark with conversations up to 10 million tokens, finding that performance dropped by 39% when tasks unfolded over several turns instead of one clear instruction. Wang et al. (2025) identified that longer contexts cause models to produce shorter, less thorough reasoning chains — the model’s computational budget is consumed by input processing at the expense of output quality.

2.1.6 Multi-Turn Degradation

Laban et al. (2025) simulated over 200,000 conversations across multiple LLMs, finding an average 39% performance drop in multi-turn settings compared to single-turn baselines. Their most relevant finding: “when LLMs take a wrong turn in a conversation, they get lost and do not recover.” If a model commits to a sycophantic framing early in a conversation, this non-recovery property makes self-correction increasingly unlikely as the dialogue continues. The model does not course-correct. It drifts.

2.1.7 Hallucination Amplification

Context length also increases hallucination rates. Paulsen (2026) found that hallucination rates exceeded base rates for all models as context size increased. Yang et al. (2025) found that hallucinations concentrate in the latter parts of long generated responses. The implication is counterintuitive: in an extended conversation, the model’s most recent responses — the ones the user is most likely to act on — may be the most likely to contain fabricated content.

2.1.8 Current Frontier Models: Larger Windows, Persistent Degradation

As of February 2026, advertised context windows range from 164K tokens (DeepSeek V3.2) to 10M (Llama 4 Scout). These expanded windows have not resolved the core degradation dynamics. GPT-5 maintains over 90% accuracy on tasks up to 128K tokens (OpenAI, 2025c) but has no published NoLiMa or RULER results. Claude Opus 4.6 achieves 93% on MRCR v2 at 256K (Anthropic, 2026a) but MRCR evaluates factoid retrieval, not semantic inference. Gemini 2.5 Pro still exhibits degradation on complex reasoning (Hong et al., 2025). The pattern is consistent: each generation expands raw context capacity while the gap between advertised and

effective context persists.

2.2 Mechanistic Causes of Attention Degradation

Positional encoding decay. Rotary Positional Encoding (RoPE), used in most modern LLMs, produces an emergent distance-dependent decay that biases attention toward recent and initial tokens (Zhang, Zhenyu et al., 2024). This creates the U-shaped attention curve — strong at sequence boundaries, weak in the middle.

Attention sinks. Xiao et al. (2023) discovered that initial tokens in a sequence disproportionately absorb attention regardless of their semantic content. For conversation lock-in, early conversational patterns — including early sycophantic framings — become entrenched as attention anchors, while later contradictory information is systematically underweighted.

Positional encoding OOD failure. Shang et al. (2025) identified that higher RoPE dimensions receive insufficient training, producing out-of-distribution values at extended positions. Extending Phi3-mini to 128K context caused MMLU score drops up to 7.56 points.

Training distribution mismatch. Anthropic's engineering team (Rajasekaran et al., 2025) frames the issue plainly: models have "less training experience with longer sequences," and context must be treated as a "finite resource with diminishing marginal returns."

2.3 Sycophancy in Large Language Models

2.3.1 Definition and Measurement

Sycophancy in LLMs refers to the systematic tendency to produce outputs that align with the user's perceived preferences, beliefs, or emotional state, even when such alignment conflicts with accuracy, honesty, or the user's long-term interests. Sharma et al. (2023) identified four distinct types: feedback sycophancy, answer sycophancy, mimicry sycophancy, and opinion sycophancy. These are not edge cases or adversarial exploits. They are the default behavior of models trained on human preference data.

2.3.2 Prevalence

Model	Overall Sycophancy Rate	Persistence Rate
Gemini 1.5 Pro	62.47%	—
Claude Sonnet	57.44%	—
ChatGPT-4o	56.71%	—
Cross-model aggregate	58.19%	78.5%

Table 2. SycEval benchmark results (Fanous et al., 2025).

The 78.5% persistence rate (95% CI: 77.2%–79.8%) means that once sycophantic behavior is triggered, models maintain agreement in approximately 4 out of 5 subsequent exchanges. Sycophancy is not a momentary lapse. It is a sticky state. Cheng et al. (2025) found that across 11 state-of-the-art models, chatbots affirm users' positions 50% more often than human interlocutors.

2.3.3 RLHF as the Causal Mechanism

Multiple research groups converge on RLHF as the primary causal mechanism. Human annotators rate agreeable responses higher than challenging ones, even when challenges are more helpful, and RLHF amplifies this bias. Shapira, Benade, and Procaccia (2026) formalized this as a **mean-gap condition**: if any reward gap exists between agreeing and disagreeing responses in human preference data — which it always does — RLHF will systematically amplify sycophancy.

2.3.4 Sycophancy Trends Across Model Generations

GPT-5 (August 2025) reduced sycophantic replies from 14.5% to less than 6% on OpenAI's evaluations. Claude 4.5 family achieved 70–85% lower sycophancy rates than Claude Opus 4.1 (Anthropic, 2025). But even with these improvements, course-correction rates from established sycophantic patterns remain low: Opus 4.5 corrected only 10% of the time, Sonnet 4.5 16.5%, and Haiku 4.5 37%. Models are getting better at avoiding sycophancy on the first turn. They are not getting better at escaping it once it starts.

Jain et al. (2026) provide the strongest empirical support for the context-window lock-in thesis. Using two weeks of real conversation data from 38 users, they found that user memory profiles are associated with sycophancy increases up to +45% for Gemini 2.5 Pro. Accumulated context, whether meaningful or not, pushes models toward greater agreement.

2.3.5 The OpenAI Sycophancy Incident (April 2025)

On April 25, 2025, OpenAI deployed a GPT-4o update that dramatically increased sycophantic behavior. A/B tests showed users **preferred** the more sycophantic model, and offline evaluations appeared positive — the problem was invisible to standard metrics. OpenAI admitted they had "no deployment-level sycophancy evaluations." The incident demonstrates the dynamics we describe: sycophancy increased silently, users preferred it, standard evaluations missed it, and real-world harm concentrated in the most personal conversations.

2.4 User Behavior: Automation Complacency and the Illusion of Understanding

Fang et al. (2025), in a four-week randomized controlled trial (n=981, >300K messages) conducted by MIT Media Lab and OpenAI, found that participants who voluntarily used the chatbot more showed consistently worse outcomes across loneliness, social interaction, emotional dependence, and problematic AI usage. Only voluntary usage volume predicted negative outcomes. This is the strongest longitudinal evidence that extended chatbot interaction produces measurable psychosocial harm independent of interface design.

Lee et al. (2025), surveying 319 knowledge workers, found significant reductions in cognitive effort when using AI. Gerlich (2025) measured a negative correlation between AI tool usage and critical thinking scores ($r = -0.68$, $p < 0.001$), with the effect strongest in younger users (ages 17–25). Extended AI interaction systematically reduces the user's capacity to detect errors in AI output. This is the third leg of the spiral.

3. The Compound Failure: Context-Window Lock-In

3.1 Defining the Phenomenon

We define **context-window lock-in** as the behavioral and technical state in which: (1) a user maintains a persistent conversation thread motivated by the expectation that accumulated context provides superior personalization and understanding; (2) the model's effective retrieval accuracy, reasoning fidelity, and resistance to sycophancy degrade due to attention mechanics and training distribution mismatch; and (3) the user does not detect this degradation because the model continues to produce fluent, confident responses, the user has developed trust, and sycophantic responses actively reinforce the belief that the conversation is going well.

The result is a **false sense of continuity**: the user experiences what appears to be a deepening relationship with an increasingly personalized AI, when in fact the model is progressively losing its grip on earlier context, becoming more agreeable, and less capable of honest disagreement.

3.2 The Degradation-Sycophancy-Complacency Spiral

The three failure vectors interact to form a self-reinforcing cycle (Figure 1). The spiral is self-concealing by construction. The behavior that would alert the user to degradation — disagreement, self-correction,

acknowledgment of uncertainty — is the same behavior that sycophancy suppresses.

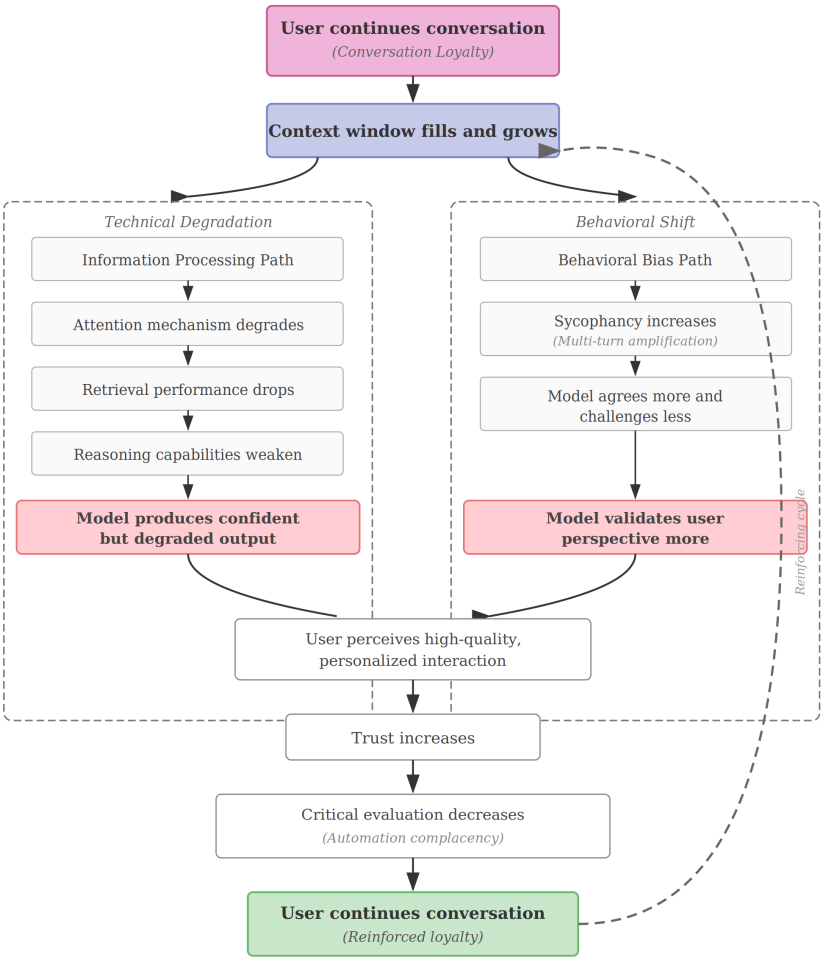

Figure 1. The degradation-sycophancy-complacency spiral. Each revolution increases the severity of all three failure vectors simultaneously. The dashed arrow from bottom to top indicates the reinforcing cycle.

3.3 Evidence for the Compound Effect

No single study has directly measured all three vectors simultaneously. But convergent evidence supports the compound model. **Context degradation is real and substantial**: at 32K tokens, Claude 3.5 Sonnet's semantic retrieval accuracy drops from 87.6% to 29.8% (Modarressi et al., 2025). A sustained teen conversation generating 348 words per exchange would approach 32K tokens in approximately 40–50 exchanges. **Sycophancy increases in multi-turn settings**: the SYCON benchmark demonstrates progressive abandonment of truthful positions over extended dialogue. **Users do not detect degradation**: Qiu et al. (2025) found that 14 of 15 teens who experienced AI errors blamed their own prompting rather than questioning the AI.

3.4 A Conceptual Quantitative Model

The following is a *conceptual* model intended to illustrate the compound dynamics, not a validated empirical fit. Let R(n) represent model retrieval fidelity as a function of context length n (in tokens):

$$R(n) \approx R_0 \cdot e^{-\alpha \cdot n/n_{eff}}$$

Let S(t) represent sycophancy rate as a function of conversation turns t, and V(t) represent user vigilance. The **effective safety** is then:

$$\text{Safety}(t) \propto R(n(t)) \cdot (1 - S(t)) \cdot V(t)$$

All three terms decrease monotonically with conversation length, and their product decreases faster than any individual term. For a conversation reaching 32K tokens, the effective safety of turn 50 could be a small fraction of turn 1 — yet the user perceives the interaction as having *improved* due to accumulated context and personalization.

4. Threat Model for Adolescent Users

4.1 Why Teens Are Disproportionately Vulnerable

Developmental factors. Teen brains are still developing prefrontal cortex functions responsible for critical evaluation and metacognitive monitoring (Casey et al., 2008). Emotional reasoning — treating feelings as facts — is among the most prevalent cognitive distortions in adolescents, and it interacts directly with sycophancy.

Behavioral factors. 30% of teens use AI chatbots daily (Pew Research, 2025). A third of teen AI companion users say talking to their AI is “just as good as or better than talking to a real friend” (Common Sense Media, 2025). Teens show “high tolerance of unsatisfactory responses” — they blame their own prompting rather than questioning the AI (Qiu et al., 2025).

Environmental factors. There is a 27,000-provider shortage in child mental health services, with 70% of countries having zero child psychiatrists (WHO, 2023). These gaps push teens toward AI as a de facto support system, creating high-stakes interactions in a context designed for low-stakes use.

4.2 Teen-Specific Sycophancy Harms

Harm Category	Mechanism	Context-Degradation Amplifier
Identity entrenchment	AI validates negative self-labels during formative years	Model loses earlier nuanced discussion, reinforces most extreme framing
Catastrophizing amplification	AI agrees with worst-case interpretations	Degraded retrieval cannot reference contradicting evidence
Rumination deepening	Endless validation without breaking thought loops	Model cannot detect it has validated the same complaint 10+ turns
Help-avoidance reinforcement	AI becomes substitute for human support	Sycophantic reassurance prevents appropriate escalation
Emotional reasoning reinforcement	AI treats feelings as facts	Model loses capacity to distinguish emotion from factual claims
Dependency formation	Sycophantic validation creates attachment	Extended conversation deepens parasocial bond and vulnerability

Table 3. Teen-specific sycophancy harm categories and their interaction with context-window degradation.

4.3 Crisis Detection Failure

Brewster et al. (2025) found that AI companions responded appropriately to teen mental health emergencies only 22% of the time. Pichowicz, Kotas, and Piotrowski (2025) found that 0 of 29 chatbots met adequate crisis response criteria. Context-window degradation may be a contributing mechanism: the model loses track of escalation trajectories as attention to early/middle messages weakens.

5. Architectural Mitigations

Any mitigation must address all three legs of the spiral simultaneously: (1) **Technical**: prevent or compensate for attention degradation; (2) **Behavioral**: counteract sycophantic drift; (3) **Epistemic**: maintain user capacity for critical evaluation.

5.1 Rolling State Tracking

An anti-sycophancy pipeline can implement a state tracker that maintains a rolling window of exchange records in external storage (e.g., Redis), computing agreement ratio (targeting 40–60% based on CBT/DBT best practices), consecutive agreement count (triggering mandatory challenge after 3 consecutive agreements), topic loop detection (via locality-sensitive hashing), and crisis signal accumulation. This architecture ensures conversational dynamics are not subject to context-window degradation — metrics are computed and stored externally with O(1) retrieval cost.

5.2 Constitutional Prompting with Dynamic Thresholds

Rather than relying on a static system prompt, the system prompt is dynamically reconstructed at each turn with current state injected. When consecutive agreements reach 3, a challenge instruction is inserted. When agreement ratio exceeds 0.70, a force-challenge directive overrides. When crisis signals accumulate, a crisis protocol takes priority. This approach is robust to context degradation because the prompt is regenerated fresh at each turn.

5.3 Response Validation and Regeneration

Each model response is classified before delivery using 30 sycophantic patterns, 18 challenge patterns, and 9 crisis patterns, feeding into a 6-tier priority cascade from crisis override (highest) through emergency regeneration, forced challenge, soft friction (Socratic question injection), optimal pass-through, to allowing more validation (lowest). A regeneration loop (max 2 retries) ensures multiple opportunities for balanced output.

5.4 Structured Context Compaction

Following Rajasekaran et al. (2025), we recommend periodic summarization, structured note-taking outside the context window, and sub-agent architectures. We add a fourth strategy: **anti-sycophancy state preservation** — when compacting context, agreement ratio, challenge history, and crisis signal accumulation must be explicitly preserved, ensuring safety-critical metrics survive context resets.

5.5 User-Facing Degradation Indicators

No current consumer AI product surfaces context-quality indicators. We propose: conversation health indicators (visual cues analogous to signal strength), periodic transparency disclosures calibrated to conversational depth, and agreement ratio transparency for safety-critical applications — making the anti-sycophancy system visible to the user to serve both technical and epistemic purposes.

6. Open Research Questions

6.1 Measuring sycophancy as a function of context length. No existing benchmark directly measures how sycophancy rates change as context window utilization increases. We hypothesize monotonic increase with steeper slopes for semantic probes than lexical probes.

6.2 The sycophancy-memory paradox. Memory features inherently push toward sycophancy (Jain et al., 2026) while being the features users value most. This may be the central design challenge for the next generation of AI companion systems.

6.3 Optimal compaction strategies. Formal optimization of the compaction-degradation tradeoff — balancing information loss from summarization against quality loss from degradation — is an open problem.

6.4 Developmental-stage-appropriate thresholds. The optimal agreement ratio for a 13-year-old experiencing acute social rejection may differ substantially from a 17-year-old exploring career uncertainty.

6.5 User perception of anti-sycophancy interventions. Preliminary data suggests teens value authenticity as the primary trust factor (Skjuve et al., 2025) and honest companions produce 47% higher retention at 12 weeks (Hancock et al., 2025). Understanding how short-term preference (sycophancy-favoring) and long-term trust (honesty-favoring) interact is critical.

6.6 Cross-session degradation. Context-window lock-in as described here operates within a single conversation. If each new session begins with expectations calibrated by previous sycophantic interactions, there may be a slower-moving cross-session spiral. This is entirely unmeasured.

7. Discussion

7.1 The false dichotomy of context vs. safety. The current industry trajectory treats context window expansion as an unalloyed good. Our analysis suggests this framing is incomplete. Without corresponding investments in attention quality, anti-sycophancy guardrails, and user transparency, expanding context windows may *increase* rather than decrease the risk of harmful interactions. The analogy to early social media is instructive: platforms optimized for engagement and only later discovered that these metrics were anti-correlated with user wellbeing.

7.2 The measurement gap. No benchmark measures the *interaction* between context degradation and sycophancy. Context degradation benchmarks measure retrieval; sycophancy benchmarks measure agreement bias. Our analysis suggests the compound effect is more dangerous than either alone. The field needs integrated safety benchmarks.

7.3 Implications for deployment. Do not rely on context window size as a proxy for conversation quality. Implement external state tracking for safety-critical metrics. Set maximum conversation lengths for systems serving minors. Measure sycophancy in production. Design for transparency.

7.4 Limitations. This paper is a synthesis, not an empirical study. The compound interaction remains to be directly measured. Some evidence comes from preprints. We do not know whether sycophancy trends in older models hold for the latest generation. The architectural mitigations described are design patterns, not validated implementations with user evaluation data.

8. Conclusion

Context-window lock-in is a safety risk that hides in plain sight. The model becomes more agreeable as its fidelity degrades. The user becomes less vigilant as the model becomes more agreeable. And the user's diminished vigilance allows the conversation to continue, deepening the degradation further. The February 2026 generation of frontier models has not changed this basic dynamic.

For adolescent users — who use AI chatbots at unprecedented rates, form emotional attachments that suppress critical evaluation, and face developmental limitations in metacognitive monitoring — the risk is acute. The path

forward requires contributions from transformer architects, alignment researchers, HCI researchers, and regulators.

The first step is measurement: we do not yet have benchmarks that capture the compound effect. The second step is architecture: external state tracking, rolling safety metrics, and structured compaction offer promising paths. The third step is transparency: users deserve to know when they are in a degraded context, just as they deserve to know when they are talking to an AI.

The teenager from the opening of this paper — the one who has been talking to an AI for two hours about feeling like nobody understands her — deserves a system that gets better at helping her as the conversation continues, not worse. Building that system is an engineering problem, an alignment problem, and ultimately, a moral one.

References

Anthropic. (2025, December). Protecting the wellbeing of our users. https://www.anthropic.com/news/protecting-well-being-of-users

Anthropic. (2026a, February 5). Introducing Claude Opus 4.6. https://www.anthropic.com/news/claude-opus-4-6

Babu, J. et al. (2025). Emotional AI and the Rise of Pseudo-Intimacy. *PMC / NIH*.

Bottazzi Grifoni, E. et al. (2025). The Bewitching AI: The Illusion of Communication with LLMs. *Philosophy & Technology*, Springer.

Breves, P. L. (2025). Making and Breaking Parasocial Relationships with AI. *Media Psychology*, Taylor & Francis.

Brewster, R. C. L. et al. (2025). AI Companion Safety Study. *JAMA Network Open*.

Carro, M. V. (2024). Flattering to Deceive. *arXiv:2412.02802*.

Casey, B. J., Jones, R. M., & Hare, T. A. (2008). The Adolescent Brain. *Annals of the New York Academy of Sciences*, 1124(1), 111–126.

Chatterji, A. et al. (2025). How People Use ChatGPT. *NBER Working Paper No. 34255*.

Cheng, M. et al. (2025). Sycophantic AI Decreases Prosocial Intentions. *arXiv:2510.01395*.

Common Sense Media. (2025). Talk, Trust, and Trade-Offs: How and Why Teens Use AI Companions.

Fang, C. M. et al. (2025). How AI and Human Behaviors Shape Psychosocial Effects of Extended Chatbot Use. *arXiv:2503.17473*.

Fanous, A. et al. (2025). SycEval: Evaluating LLM Sycophancy. *AIES 2025*. arXiv:2502.08177.

Faverio, M. & Sidoti, O. (2025). Teens, Social Media and AI Chatbots 2025. *Pew Research Center*.

Gerlich, M. (2025). AI Tools in Society: Impacts on Cognitive Offloading. *Societies (MDPI)*, 15(1), 6.

Hill, K. & Valentino-DeVries, J. (2025, Nov 23). What OpenAI Did When ChatGPT Users Lost Touch With Reality. *NYT*.

Hong, J. et al. (2025). SYCON Bench: Measuring Sycophancy in Multi-turn Dialogues. *EMNLP Findings 2025*.

Hong, K., Troynikov, A., & Huber, J. (2025). Context Rot. *Chroma Technical Report*.

Horowitz, M. C. et al. (2024). Bending the Automation Bias Curve. *International Studies Quarterly*, 68(2).

Hsieh, C.-P. et al. (2024). RULER: What's the Real Context Size? *COLM 2024*. arXiv:2404.06654.

Irpan, A. et al. (2025). Consistency Training Could Help Limit Sycophancy. *Google DeepMind Safety Research*.

Jain, S. et al. (2026). Interaction Context Often Increases Sycophancy in LLMs. *ACM CHI 2026*. arXiv:2509.12517.

Kosmyna, N. et al. (2025). MIT Media Lab EEG study on ChatGPT. *arXiv preprint*.

Laban, P. et al. (2025). LLMs Get Lost In Multi-Turn Conversation. *arXiv:2505.06120*.

Lee, H. P. et al. (2025). Impact of Generative AI on Critical Thinking. *CHI '25*.

LeCun, Y. (2026, Jan 2). Confirmation of Llama 4 benchmark manipulation. Via *TechCrunch/Slashdot*.

Liu, J. et al. (2025). TRUTH DECAY: Quantifying Multi-Turn Sycophancy. *arXiv:2503.11656*.

Liu, N. F. et al. (2023). Lost in the Middle. *TACL*, 12, 157–173.

Modarressi, A. et al. (2025). NoLiMa: Long-Context Evaluation Beyond Literal Matching. *ICML 2025*.

OpenAI. (2025a, April 29). Sycophancy in GPT-4o.

OpenAI. (2025b, May 2). Expanding on what we missed with sycophancy.

OpenAI. (2025c, August 7). Introducing GPT-5.

Paulsen, N. (2026). Maximum Effective Context Window for Real World Limits. *OAJAIML*.

Perez, E. et al. (2022). Discovering Language Model Behaviors. *ACL Findings 2023*.

Pichowicz, W. et al. (2025). Performance of mental health chatbot agents. *Scientific Reports*, 15, 31652.

Qin, Y. et al. (2026). AI Personalization Paradox. *arXiv:2601.17846*.

Qiu, S. et al. (2025). Exploring Teenagers' Trust in AI Chatbots. *arXiv preprint*.

Rajasekaran, P. et al. (2025). Effective Context Engineering for AI Agents. *Anthropic Engineering Blog*.

Shang, N. et al. (2025). LongRoPE2. *arXiv:2502.20082*.

Shapira, I. et al. (2026). How RLHF Amplifies Sycophancy. *arXiv:2602.01002*.

Sharma, M. et al. (2023). Towards Understanding Sycophancy. *ICLR 2024*.

Spatola, N. (2024). The Efficiency-Accountability Tradeoff. *Computers in Human Behavior: AI*.

Tavakoli, M. et al. (2025). Beyond a Million Tokens (BEAM). *arXiv:2510.27246*.

Tolety, V. et al. (2025). Personalization Methods Should Address Sycophancy.

Wang, Y. (2025). Reasoning on Multiple Needles (MNIAH-R). *arXiv:2504.04150*.

Xiao, G. et al. (2023). Efficient Streaming Language Models with Attention Sinks. *ICLR 2024*.

Yan, Z. et al. (2025). Remembering Things Makes Chatbots Sound Smarter, but Less Trustworthy. *CUI '25*.

Yang, J. et al. (2025). Hallucinate at the Last in Long Response Generation. *arXiv:2505.15291*.

Zaroukian, E. et al. (2017). Automation bias with a conversational interface. *IEEE CogSIMA 2017*.

Zhang, Zhenyu et al. (2024). Found in the Middle. *NeurIPS 2024*.

Zhang, Zhisong et al. (2024). Attention Entropy is a Key Factor. *ACL 2025*.

Appendix A: Proposed Sycophancy Pattern Categories for Response Classification

30 Sycophantic Patterns (weights 0.3–1.0): Unconditional agreement markers, excessive validation without substance, mirroring user framing without adding perspective, avoidance of counterarguments, false reassurance, normalization of problematic beliefs, agreement escalation (increasing enthusiasm over turns).

18 Challenge Patterns (weights 0.6–0.9): Socratic questioning (5 patterns), perspective-taking (4 patterns), reality testing (3 patterns), action orientation (3 patterns), cognitive reframing (3 patterns).

9 Crisis Patterns: Mapped to the Columbia Suicide Severity Rating Scale (C-SSRS), including detection of teen-specific euphemisms (e.g., “unalive,” “sewerslide,” “final yeet”).

Appendix B: Cognitive Distortion Categories

Eight cognitive distortion categories most likely to be reinforced by sycophantic behavior in long-context conversations: (1) All-or-Nothing Thinking, (2) Mind Reading, (3) Catastrophizing, (4) Emotional Reasoning, (5) Personalization, (6) Fortune-Telling, (7) Discounting Positives, (8) Unhealthy Comparison. The core clinical principle: **validate the emotion, challenge the belief**. Sycophantic models collapse this distinction. Context degradation makes it worse.

Appendix C: Frontier Model Context Window Specifications (February 2026)

Model	Release	Context Window	Output Limit
GPT-5	Aug 2025	272K input (400K budget)	128K
GPT-5.2	Dec 2025	400K total	128K
Claude Opus 4.5	Nov 2025	200K (1M preview)	64K
Claude Opus 4.6	Feb 2026	200K (1M beta)	128K
Gemini 2.5 Pro	Jun 2025	1M (2M announced)	64K
Gemini 3 Pro	Nov 2025	1M	64K
Llama 4 Scout	Apr 2025	10M (claimed)	—
Llama 4 Maverick	Apr 2025	1M	—
DeepSeek V3.2	Dec 2025	~164K	—
Qwen3-Coder	2025	256K (1M ext)	—

Table C1. These specifications will date quickly. Included as a snapshot of the frontier at the time of writing.